

Empirical Investigation of Multi-Sensor Fusion, Representation Collapse, and Policy Stability in CARLA Reinforcement Learning

Bachelor Thesis Project (BTP)
Department of Computer Science and Engineering
Indian Institute of Technology
Project Codebase: <https://github.com/MSR327/vlm>

Abstract—Deep reinforcement learning (RL) for vision-based autonomous driving requires balancing perceptual coverage, state dimensionality, and policy optimization stability. We present a systematic empirical study on sensor scaling and representation learning in the CARLA simulator using Proximal Policy Optimization (PPO). Rather than reporting only an idealized endpoint, we evaluate an empirical ladder across four sensory configurations: (1) a monocular baseline (100-dim) achieving 39.0% route completion (195 m, $R = 313.21 \pm 38.45$, lane deviation 0.42 ± 0.11 m) in Town01 but failing at sharp intersections (Meter 190) and degrading under zero-shot transfer in Town02 ($14.2\% \pm 3.1\%$ completion); (2) a multi-view forward surround rig with 3 cameras (290-dim, $0^\circ, \pm 60^\circ$) reaching 25.0% route completion (125 m, $R = 158.31 \pm 26.80$, lane deviation 0.28 ± 0.08 m), a degradation from the monocular baseline despite wider coverage; (3) a naive 4-camera panoramic surround rig (385-dim) that induces policy standstill collapse (3.4% completion, 17 m, $R = 17.35 \pm 3.12$, lane deviation 0.05 ± 0.02 m) due to high-dimensional gradient dissipation and reward-design exploitation; and (4) an orthogonal multimodal fusion rig (195-dim) pairing front vision with a 2D Bird’s-Eye-View (BEV) LiDAR height grid, achieving reward improvement ($-8.81 \rightarrow -5.02 \pm 4.18$, lane deviation 0.35 ± 0.09 m) with zero stalls but only 6.0% route completion. Furthermore, we isolate and resolve two severe numerical instabilities in sparse BEV autoencoding: a float32 exponential overflow (NaN loss) cured via bounded log-sigma clipping $[-15.0, 1.0]$, and a 1.8×10^6 latent variance explosion cured via unconditional forward execution. Finally, we analyze the structural limits of static latent concatenation across independent VAEs, observing that ungrounded MLP policy heads struggle to reconcile high-dimensional multi-modal features, motivating future investigation of patch-based cross-attention and Vision-Language-Action (VLA) architectures.

Index Terms—Autonomous driving, CARLA simulator, sensor fusion, Bird’s-Eye-View LiDAR, Proximal Policy Optimization, Variational Autoencoder, representation learning, cross-attention.

I. INTRODUCTION & MOTIVATION FOR MULTI-SENSOR SCALING

End-to-end autonomous driving aims to map onboard perceptual sensor streams directly to continuous steering, throttle, and brake commands [1]–[3]. While foundational imitation learning and reinforcement learning agents operated primarily on a single forward-facing RGB camera, navigating real-world or simulated urban environments demands wider spatial perception and reliable 3D distance awareness [4], [6].

A. Engineering Motivation for Progressive Sensor Scaling

Deploying autonomous agents in urban environments exposes four fundamental physical and geometric constraints of single-sensor architectures:

- 1) **Geometric Blind Spots at Urban Intersections:** Standard monocular perspective cameras provide a horizontal field of view (FOV) of 90° – 125° . While sufficient for straight highway tracking, this geometry creates complete perceptual blindness to lateral corridors ($> 60^\circ$ off-center). At sharp 90° urban turns or T-junctions, the inner lane boundary and road curvature exit the front camera frame before the vehicle can physically rotate its yaw axis, leaving the control policy without guidance.
- 2) **Ill-Posed Inverse Depth Ambiguity:** Projecting 3D physical coordinates (x, y, z) onto a 2D pixel array (u, v) discards metric range. Monocular depth estimation is fundamentally ill-posed and susceptible to scale ambiguity, dynamic illumination shifts, rain spray, and sun glare [16].
- 3) **Occlusion Resilience in Dense Traffic:** In dense traffic, leading trucks, buses, or street furniture occlude the forward path. Multi-view and surround camera arrays provide parallax and line-of-sight redundancy around obstacles.
- 4) **Photometric vs. Geometric Complementarity:** While vision provides dense semantic classification (lane markings, traffic signs, stop lines), active LiDAR sensors provide metric, illumination-invariant spatial ranges with centimeter precision. Fusing camera vision with LiDAR Bird’s-Eye-View (BEV) representations combines semantic context with absolute metric geometry.

However, simply adding more sensors naively into reinforcement learning pipelines frequently degrades performance rather than improving it. In this work, we construct and evaluate an empirical four-stage *sensor ladder* in CARLA to investigate the trade-offs between perceptual coverage, state dimensionality, simulation stability, and policy convergence.

II. RELATED WORK

A. Sensor Fusion in Autonomous Driving

Learning by Cheating (LBC) [6] demonstrated that agents trained on privileged Bird’s-Eye-View (BEV) abstractions substantially outperform monocular camera agents, highlighting the performance gap between perspective-view and metric spatial representations. *TransFuser* [5] demonstrated that combining multi-view cameras with LiDAR point clouds via cross-attention outperforms single-modality baselines by fusing optical semantic segmentation with invariant metric 3D geometry. *InterFuser* [7] extended multi-sensor fusion across panoramic perspectives to eliminate blind spots and ensure collision safety in complex traffic scenarios.

B. Bird’s-Eye-View (BEV) Point Cloud Representations

Projecting raw 3D LiDAR point clouds into 2D Bird’s-Eye-View (BEV) spatial grids preserves metric physical distances while enabling standard 2D convolutional backbones to extract geometric affordances efficiently: *PointPillars* [8] pioneered vertical pillar tokenization for real-time 3D object detection from point clouds. *BEVDet* [9] further established that transforming perceptual features into an ego-centric 2D metric ground plane provides superior spatial localization and trajectory prediction compared to perspective-view processing.

C. Representation Learning & Dimensionality in RL

Model-free reinforcement learning directly from raw high-dimensional pixels suffers from extreme sample inefficiency and optimization instability [10], [12]. *CARL* [10] demonstrated in CARLA benchmark evaluations that policy gradient convergence is heavily constrained by observation representation dimensionality. *Roach* [11] addressed high-dimensional sample inefficiency by pre-training policies over privileged BEV affordance masks. Variational Autoencoders (VAEs) [15] are widely deployed to compress high-dimensional observation frames into compact, regularized continuous latent vectors $\mathbf{z} \in \mathbb{R}^{d_z}$ to feed downstream policy heads. Furthermore, abstracting natural RGB streams into Cityscapes-style semantic class palettes [16] reduces perceptual complexity. In simulation, semantic segmentation masks are provided as ground-truth labels directly by the rendering engine, bypassing the need for learned perception.

D. Policy Optimization & Control Stability

Proximal Policy Optimization (PPO) [12] represents the primary on-policy clipped surrogate reinforcement learning objective for continuous control. Generalized Advantage Estimation (GAE- λ) [13] introduces exponential discounting to calibrate the bias-variance trade-off in policy advantage estimates. Conditioning for Action Policy Smoothness (CAPS) [14] demonstrated that model-free continuous RL agents frequently suffer from high-frequency actuator chatter unless explicitly regularized via action smoothness loss penalties.

III. METHODOLOGY

A. Reinforcement Learning Formulation

We formulate the closed-loop autonomous driving problem as an episodic Partially Observable Markov Decision Process (POMDP) defined by the tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$:

- **Control Frequency:** The agent interacts synchronously with the CARLA simulator at a fixed frequency of $f = 20$ Hz ($\Delta t = 0.05$ s).
- **Action Space:** The policy outputs a continuous 2-dimensional action vector $\mathbf{a}_t = [a_{\text{steer}}, a_{\text{long}}] \in [-1, 1]^2$.
- **Decoupled Longitudinal Control:** To prevent simultaneous, conflicting throttle and brake commands, we implement a signed longitudinal mapping:

$$\text{throttle}_t = \max(a_{\text{long}}, 0.0) \quad (1)$$

$$\text{brake}_t = \max(-a_{\text{long}}, 0.0) \quad (2)$$

- **Temporal Action Smoothing:** Guided by the CAPS framework [14], we filter steering predictions using an Exponential Moving Average (EMA) to prevent vehicle instability and actuator chatter:

$$\delta_t = 0.6 \delta_{t-1} + 0.4 a_{\text{steer},t} \quad (3)$$

where δ_t denotes the physical steering angle applied to the CARLA ego-vehicle.

- **Multi-Objective Reward Function:** At each discrete time step t , the agent receives a shaped reward:

$$r_t = R_{\text{speed}} \cdot R_{\text{center}} \cdot R_{\text{angle}} + R_{\text{penalty}} \quad (4)$$

where the individual components are defined as:

$$R_{\text{speed}} = \min(v_t/22.0, 1.0) \quad (5)$$

$$R_{\text{center}} = \max\left(1.0 - \frac{|d_{\text{center}}|}{3.0}, 0.0\right) \quad (6)$$

$$R_{\text{angle}} = \max\left(1.0 - \frac{|\theta_{\text{heading}}|}{20^\circ}, 0.0\right) \quad (7)$$

Here v_t represents scalar forward speed in km/h, d_{center} denotes signed perpendicular lane-center deviation in meters, and θ_{heading} represents the angular heading error relative to the lane tangent. Dynamic collisions, sidewalk encroachments, or lane deviations exceeding 3.0 m terminate the episode with a penalty of $R_{\text{penalty}} = -10.0$.

B. The Empirical Sensor Ladder

To isolate the empirical trade-offs between perceptual coverage, representation dimensionality, and policy stability, we evaluate four systematic sensor rungs (Fig. 1):

- 1) **Rung 1: Monocular Baseline (100-dim):** A single front-facing semantic camera (0° yaw, 125° FOV). The state vector concatenates the 95-dimensional visual latent with a 5-dimensional navigation telemetry vector:

$$\mathbf{s}_t = [\mathbf{z}_{\text{cam}} \in \mathbb{R}^{95} \parallel \mathbf{x}_{\text{nav}} \in \mathbb{R}^5] \in \mathbb{R}^{100} \quad (8)$$

where $\mathbf{x}_{\text{nav}} = [v_t, d_{\text{center}}, \theta_{\text{heading}}, \delta_{t-1}, a_{\text{long},t-1}]$.




Fig. 1. Architectural diagrams of the four sensor ladder configurations evaluated in CARLA. (a) Rung 1: Monocular front baseline. (b) Rung 2: Multi-view forward surround rig with 3 cameras (210° panoramic forward arc). (c) Rung 3: Naive 4-camera panoramic surround rig (360° horizontal coverage). (d) Rung 4: Orthogonal multimodal fusion of front camera and 2D BEV LiDAR.

- 2) **Rung 2: Multi-View Forward Surround (290-dim):** To address the severe turn blindness of the monocular baseline while preserving forward motion dynamics, the rig integrates three forward-focused cameras: Front (0° yaw, 125° FOV), Left-Angled (−60° yaw, 90° FOV), and Right-Angled (+60° yaw, 90° FOV), yielding a combined forward panoramic arc of 210°:

$$s_t = [z_{\text{front}} \parallel z_{\text{left}} \parallel z_{\text{right}} \parallel x_{\text{nav}}] \in \mathbb{R}^{290} \quad (9)$$

- 3) **Rung 3: Naive 4-Camera Panoramic Surround (385-dim):** Four orthogonal semantic cameras covering the full 360° horizontal perimeter: Front (0°, 125° FOV), Left (−90°, 90° FOV), Right (+90°, 90° FOV), and Rear (180°, 90° FOV). The state is formed by direct latent concatenation:

$$s_t = [z_{\text{front}} \parallel z_{\text{left}} \parallel z_{\text{right}} \parallel z_{\text{rear}} \parallel x_{\text{nav}}] \in \mathbb{R}^{385} \quad (10)$$

- 4) **Rung 4: Orthogonal Multimodal Fusion (195-dim):** One front-facing camera (0°, 125° FOV) combined with a roof-mounted 3D LiDAR projected into a 2D Bird’s-Eye-View (BEV) height/occupancy grid:

$$s_t = [z_{\text{cam}} \in \mathbb{R}^{95} \parallel z_{\text{lidar}} \in \mathbb{R}^{95} \parallel x_{\text{nav}} \in \mathbb{R}^5] \in \mathbb{R}^{195} \quad (11)$$

C. Simulation Environment, Vehicle Physics, & Benchmark Maps

To guarantee experimental repeatability and eliminate physics artifacts, all experiments adhere to a rigorously controlled simulation environment:

- **Simulation Platform:** CARLA version 0.9.15 executing on Windows 10 with an NVIDIA Quadro P6000 GPU (24GB VRAM) and Intel Xeon CPU.
- **Synchronous Execution & Physics Substepping:** Fixed simulation time step $\Delta t = 0.05$ s (20 Hz). Server and client operate in non-real-time synchronous mode to ensure zero frame dropping. Under the PhysX engine,

substepping is configured at 100 Hz ($\delta t_{\text{sub}} = 0.01$ s, 5 substeps per step) to guarantee rigid-body stability and consistent tire contact forces.

- **Ego-Vehicle Dynamics:** Standard four-wheeled passenger vehicle asset (Tesla Model 3 / Audi TT) with curb mass $m = 1500$ kg, wheelbase $L = 2.87$ m, and dry asphalt surface tire friction coefficient $\mu = 1.0$.
- **Sensor Mounting Coordinates (Extrinsics):**
 - *Front Camera:* Mounted on the windshield header at $(x = 2.0 \text{ m}, y = 0.0 \text{ m}, z = 1.4 \text{ m})$, orientation (pitch = −5°, yaw = 0°, roll = 0°), 125° FOV.
 - *Left-Angled Camera:* $(x = 2.0 \text{ m}, y = -0.2 \text{ m}, z = 1.4 \text{ m})$, orientation (pitch = −5°, yaw = −60°, roll = 0°), 90° FOV.
 - *Right-Angled Camera:* $(x = 2.0 \text{ m}, y = +0.2 \text{ m}, z = 1.4 \text{ m})$, orientation (pitch = −5°, yaw = +60°, roll = 0°), 90° FOV.
 - *Rear Camera (Rung 3):* $(x = -1.5 \text{ m}, y = 0.0 \text{ m}, z = 1.4 \text{ m})$, orientation (pitch = −5°, yaw = 180°, roll = 0°), 90° FOV.
 - *3D LiDAR Sensor:* Mounted on the roof center at $(x = 0.0 \text{ m}, y = 0.0 \text{ m}, z = 2.4 \text{ m})$, configured with 32 vertical laser channels spanning −30° to +10°, rotation frequency 20 Hz, outputting 100,000 points/sec over a 50 m radius.
- **Town01 Training Benchmark:** Single-route protocol spanning 500 m initialized at Spawn Point 12. Road width is 6.0 m (3.0 m per lane), bordered by 0.15 m concrete curbs. The trajectory features a long 190 m straightaway leading directly into an unbanked, blind 90° right turn with turning radius $r \approx 12.5$ m (curvature $\kappa \approx 0.08 \text{ m}^{-1}$) at *Meter 190*.
- **Town02 Zero-Shot Transfer Benchmark:** Evaluated across 20 unseen evaluation episodes. Town02 features narrow 3.0 m lanes, high-density multi-story stone masonry facades, sharp 90° blind T-junctions, and complex building geometries ($> 1.2 \times 10^6$ polygons per frame).
- **Dynamic Traffic & Environmental Variations:** A




Fig. 2. Geometric coordinate mapping and binning discretization of the 2D Bird's-Eye-View (BEV) LiDAR projection raster (80×160 channels) centered egocentrically around the vehicle.

dataset of 10,000 multi-sensor frames was collected under 29 autonomous NPC vehicles operating via CARLA Traffic Manager (15–35 km/h) and 4 randomized weather presets (ClearNoon, HardRainNoon, WetSunset, CloudyNight).

D. 2D Bird's-Eye-View (BEV) LiDAR Projection Geometry

A 360-degree raycast LiDAR sensor is top-mounted on the ego-vehicle chassis at $(x = 0.0, y = 0.0, z = 2.4 \text{ m})$. At each time step, raw 3D point cloud returns $\mathcal{P} = \{(x_i, y_i, z_i)\}_{i=1}^N$ are projected into an asymmetric Bird's-Eye-View spatial grid (Fig. 2):

- 1) **Asymmetric Spatial Bounding:** We define an asymmetric region of interest biased longitudinally forward:

$$\mathcal{P}_{\text{filt}} = \{(x_i, y_i, z_i) \in \mathcal{P} \mid -15.0 \leq x_i \leq +35.0, -25.0 \leq y_i \leq +25.0\} \quad (12)$$

- 2) **Discretization to Pixel Grid Coordinates (u_i, v_i) :** The continuous spatial coordinates are mapped into an 80×160 discrete matrix:

$$u_i = \text{clip} \left(\left\lfloor \frac{35.0 - x_i}{50.0} \times 79 \right\rfloor, 0, 79 \right) \quad (13)$$

$$v_i = \text{clip} \left(\left\lfloor \frac{y_i + 25.0}{50.0} \times 159 \right\rfloor, 0, 159 \right) \quad (14)$$

- 3) **Normalized Height Assignment:** The elevation $z_i \in [-2.0 \text{ m}, +2.0 \text{ m}]$ is mapped to continuous unit range:

$$\tilde{z}_i = \text{clip} \left(\frac{z_i + 2.0}{4.0}, 0.0, 1.0 \right) \quad (15)$$

- 4) **Max-Pooling Channel Assignment:** Where multiple LiDAR reflections project into the same spatial pixel (u, v) , max-pooling preserves the highest obstacle return:

$$\mathbf{L}_{\text{bev}}(u, v) = \max_{i: u_i=u, v_i=v} \tilde{z}_i \quad (16)$$

yielding a single-channel raster $\mathbf{L}_{\text{bev}} \in [0, 1]^{80 \times 160}$.

E. End-to-End System Pipeline & Architecture

Fig. 3 details the end-to-end signal flow from raw perception to low-level vehicle actuation. Camera and LiDAR data streams are compressed by frozen convolutional encoders, concatenated with vehicle kinematics, and processed by PPO actor-critic networks to issue smoothed continuous controls at 20 Hz.

F. Proximal Policy Optimization with GAE- λ

The policy network $\pi_\theta(\mathbf{a}_t \mid \mathbf{s}_t)$ and state-value critic $V_\phi(\mathbf{s}_t)$ are optimized using PPO with Generalized Advantage Estimation:

- 1) **Temporal Difference Residual:**

$$\delta_t^V = r_t + \gamma V_\phi(\mathbf{s}_{t+1})(1 - d_t) - V_\phi(\mathbf{s}_t) \quad (17)$$

where $d_t \in \{0, 1\}$ indicates episodic termination.

- 2) **Generalized Advantage Estimation (GAE- λ):**

$$\hat{A}_t = \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}^V \quad (18)$$

with discount $\gamma = 0.99$ and smoothing parameter $\lambda = 0.95$.

- 3) **PPO Clipped Surrogate Objective:**

$$\mathcal{L}_{\text{PPO}}(\theta) = \hat{\mathbb{E}}_t \left[\min(\rho_t \hat{A}_t, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right] \quad (19)$$

where probability ratio $\rho_t = \frac{\pi_\theta(\mathbf{a}_t \mid \mathbf{s}_t)}{\pi_{\theta_{\text{old}}}(\mathbf{a}_t \mid \mathbf{s}_t)}$ and clipping threshold $\epsilon = 0.2$.

- 4) **Value Function Mean Squared Error Loss:**

$$\mathcal{L}_{\text{VF}}(\phi) = \hat{\mathbb{E}}_t \left[(V_\phi(\mathbf{s}_t) - \hat{R}_t)^2 \right], \quad \hat{R}_t = \frac{R_t - \mu_R}{\sigma_R + 10^{-7}} \quad (20)$$

Standardizing returns $\hat{R}_t$ guarantees numerical stability across diverse episodic durations.

PPO Hyperparameters. All runs share identical optimization settings: learning rate $\alpha = 1 \times 10^{-4}$ (Adam), batch size $B = 1$ (per-step updates), $K = 15$ PPO epochs per policy update, entropy bonus coefficient $\beta_H = 0.01$, and initial action standard deviation $\sigma_{\text{init}} = 0.20$. A single random seed (seed = 42) was used for all experiments.

Evaluation Protocol. All evaluation rollouts use a *stochastic* policy (sampling from the learned Gaussian action distribution). Reported $\pm$ values represent empirical standard deviations across evaluation episodes from a single training seed, not confidence intervals across independent training runs.

IV. IMPLEMENTATION DETAILS

A. Deep Neural Network Architectures

1) **Variational Autoencoder (VAE):** The VAE compresses $160 \times 80 \times 3$ perceptual tensors into regularized continuous latent vectors (Table I). To guarantee scientific fairness across the ladder, the LiDAR BEV autoencoder in Rung 4 adopts the exact same layer geometry, parameter capacity (13.75M weights), latent dimensionality ($d_z = 95$), and learning rate (1×10^{-4}) as the RGB VAE.




Fig. 3. End-to-end multimodal deep reinforcement learning architecture in CARLA, inspired by TransFuser [5]. Perspective RGB streams and 2D metric Bird’s-Eye-View (BEV) LiDAR grids are tokenized by frozen Conv-VAE encoders into compact latent vectors ($\mathbf{z} \in \mathbb{R}^{95}$), concatenated with kinematic vehicle telemetry, and fed to PPO actor-critic heads for continuous 20 Hz closed-loop control. The callout (right) details our bounded log-sigma stabilization preventing NaN explosions alongside the empirical failure mechanisms of naive 4-camera scaling.

TABLE I
LAYER SPECIFICATIONS OF THE VARIATIONAL AUTOENCODER

Sub-module	Layer Operator	Filter / Units	Output Shape
Encoder	Input Tensor	—	$160 \times 80 \times 3$
	Conv2D (stride 2, ReLU)	$32 \times (4 \times 4)$	$80 \times 40 \times 32$
	Conv2D (stride 2, ReLU)	$64 \times (3 \times 3)$	$40 \times 20 \times 64$
	Batch Normalization	—	$40 \times 20 \times 64$
	Conv2D (stride 2, ReLU)	$128 \times (4 \times 4)$	$20 \times 10 \times 128$
	Conv2D (stride 2, ReLU)	$256 \times (3 \times 3)$	$10 \times 5 \times 256$
	Flatten & Dense	1024 units	1024
Bottleneck	Linear Head (μ)	95 units	95
	Linear Head ($\log \sigma^2$)	95 units (clipped)	95
Decoder	Dense Projection	$1024 \rightarrow 12,800$ units	$10 \times 5 \times 256$
	Conv2DTranspose (s=2)	$128 \times (4 \times 4)$	$20 \times 10 \times 128$
	Conv2DTranspose (s=2)	$64 \times (3 \times 3)$	$40 \times 20 \times 64$
	Conv2DTranspose (s=2)	$32 \times (4 \times 4)$	$80 \times 40 \times 32$
	Conv2DTranspose (s=1)	$3 \times (3 \times 3)$, Sigm.	$160 \times 80 \times 3$

2) *PPO Actor-Critic Policy*: Both actor and critic modules operate directly upon the combined sensory-navigational state $\mathbf{s}_t$:

- *Actor Network*: Multi-layer perceptron mapping $\mathbf{s}_t \rightarrow [500, \tanh] \rightarrow [300, \tanh] \rightarrow [100, \tanh] \rightarrow [2, \tanh]$.
- *Critic Network*: Multi-layer perceptron mapping $\mathbf{s}_t \rightarrow [500, \tanh] \rightarrow [300, \tanh] \rightarrow [100, \tanh] \rightarrow [1, \text{Linear}]$.
- *Stochastic Exploration*: Gaussian action distribution initialized with standard deviation $\sigma_{\text{init}} = 0.20$.

B. Numerical Stabilization on Sparse Point Clouds

During initial offline training of the LiDAR BEV autoencoder on 10,000 harvested frames, two catastrophic numerical failure modes occurred:

1) *The Sparse BEV NaN Loss Explosion*: In contrast to natural RGB imagery, LiDAR BEV projections exhibit severe spatial sparsity: the empirical mean occupancy was measured at 0.0059 (99.41% of pixels are exact zeros). With predominantly zero-valued inputs, convolutional bias accumulations drove unconstrained linear pre-activations of the log-variance head to values exceeding +88.7. In standard IEEE 754 single-precision floating point:

$$\exp(88.7228) \rightarrow +\infty \quad (21)$$

Multiplying $+\infty$ by stochastic Gaussian noise $\epsilon \sim \mathcal{N}(0, \mathbf{I})$ produced NaN values, causing gradient backpropagation to immediately corrupt all network weights. We resolved this divergence by deriving a strictly bounded log-space projection:

$$\tilde{s} = \text{clip}(\text{Dense}_\sigma(\mathbf{h}), -15.0, 1.0), \quad \sigma = \exp(\tilde{s}) \quad (22)$$

This mathematically guarantees:

$$3.05 \times 10^{-7} \leq \sigma \leq 2.718 \quad (23)$$

Coupled with gradient norm clipping (clipnorm = 1.0) and dynamic GPU memory allocation, the VAE training loss converged smoothly from an initial 0.0358 to 0.0046 without divergence.

2) *The 1.8-Million Latent Variance Explosion*: When reloading the trained VAE for closed-loop evaluation, empirical latent diagnostics revealed a standard deviation of $\sigma_z = 1,812,414.25$. Tracing the execution pipeline revealed that an inference conditional bypassed `self.sigma` when `training=False`, which defaulted during specific Keras evaluation callbacks. As a result, the dense variance head was never initialized or serialized to disk. Upon reload, uninitialized random weights evaluated to magnitude 10^6 . Enforcing unconditional execution of both μ and σ across all passes permanently stabilized the latent standard deviation to 87.27. We note that this magnitude substantially exceeds unit variance and creates a scale mismatch with the PPO actor’s $\tanh$ activations ($\tanh(x)$ saturates for $|x| > 3$), which likely contributes to the optimization difficulties observed in Rung 4.

V. RESULTS & EMPIRICAL ANALYSIS

A. Master Comparative Statistical Evaluation

Table II provides a rigorous statistical comparison across all evaluated sensor configurations in CARLA, reporting mean values along with empirical standard deviations ($\pm\sigma$) across independent evaluation episodes.

TABLE II
MASTER DRIVING PERFORMANCE & FAILURE EVALUATION ACROSS THE SENSOR LADDER IN CARLA

Empirical Metric	Rung 1 (Monocular)	Rung 2 (3-Cam Forward)	Rung 3 (4-Cam Surround)	Rung 4 (1-Cam + LiDAR BEV)
Sensory Setup	1 Front Cam (125°)	3 Cams (0°, ±60°)	4 Cams (0°, ±90°, 180°)	1 Front Cam + 1 Top LiDAR
Perceptual Modality	Semantic RGB	Semantic RGB	Semantic RGB	Semantic RGB + 2D BEV Metric Depth
Observation Dim (s_t)	100	290	385	195
Encoder Architecture	Single Frozen VAE	Shared Frozen VAE	Shared Frozen VAE	Dual Frozen VAEs
Encoder Parameters	13,749,534	13,749,534	13,749,534	27,499,068
Actor Policy Parameters	231,102	326,102	373,602	278,602
<i>Town01 Closed-Loop Benchmark Performance (Mean ± Std Dev)</i>				
Route Completion (%) ↑	39.0% ± 2.4%	25.0% ± 4.1%	3.4% ± 0.8%	6.0% ± 1.9%
Max Trajectory Distance (m) ↑	195.0 m	125.0 m	17.1 m	30.2 m
Cumulative Reward (R) ↑	313.21 ± 38.45	158.31 ± 26.80	17.35 ± 3.12	-5.02 ± 4.18
Peak Episode Reward ($R_{\max}$) ↑	+342.10	+184.20	+21.40	+19.84
Lane Deviation (d_c, m) ↓	0.42 ± 0.11 m	0.28 ± 0.08 m	0.05 ± 0.02 m [†]	0.35 ± 0.09 m
Heading Error (θ_e, rad) ↓	0.082 ± 0.021	0.045 ± 0.014	0.012 ± 0.005 [†]	0.064 ± 0.018
<i>Failure Modes & Computational Efficiency</i>				
Primary Termination Mode	Blind Turn Curb Collision	Side Friction at Curb	Zero-Velocity Standstill Trap	Lane Boundary Departure
Standstill Stall Rate (%) ↓	0.0%	0.0%	95.0%	0.0%
Inference Compute (GMACs) ↓	1.42	4.26	5.68	2.85
VRAM Footprint (MB) ↓	1420 MB	2280 MB	3150 MB	1890 MB
Simulation Throughput (FPS) ↑	19.8	17.4	13.2	18.6
<i>Town02 Zero-Shot Transfer Benchmark (Mean ± Std Dev, 20 Evaluation Episodes)</i>				
Zero-Shot Completion (%) ↑	14.2% ± 3.1%	18.5% ± 3.8%	2.1% ± 0.5%	22.4% ± 4.2%
Zero-Shot Mean Reward ↑	42.10 ± 11.2	68.40 ± 14.5	8.10 ± 2.4	91.30 ± 16.8

[†] Artificially suppressed metric: the vehicle remained motionless at distance 17.1 m in the center of the lane.




Fig. 4. Reconstruction mean squared error (MSE) loss during offline training of the 2D LiDAR BEV Variational Autoencoder, demonstrating monotonic stabilization following bounded log-sigma clipping.

B. Empirical Learning Curves & Convergence

Fig. 6 and Fig. 7 illustrate the closed-loop training performance across the four sensor ladder rungs in Town01. Rung 1 trained for $\sim 105k$ steps, achieving $39.0\% \pm 2.4\%$ route completion ($R = 313.21 \pm 38.45$). Rung 2 reached $25.0\% \pm 4.1\%$ route completion ($R = 158.31 \pm 26.80$) in 12.2k steps—a lower completion than the monocular baseline, though the substantially shorter training budget makes direct comparison difficult. Rung 3 flatlined at $3.4\% \pm 0.8\%$ completion ($R = 17.35 \pm 3.12$) in a standstill trap. Rung 4 achieved -5.02 ± 4.18 reward across 230 episodes (13.2k steps) with zero stall terminations but only 6.0% route completion. **Important caveat:** Rungs 2–4 trained for 8–10 $\times$ fewer steps than Rung 1, so performance differences may reflect training duration rather than architectural limitations. Fig. 4 confirms




Fig. 5. Distribution of driving infractions and episodic termination modes across 20 closed-loop evaluation episodes in Town01, inspired by the TransFuser benchmark protocol [5]. While Rung 1 predominantly terminates on outer curb collisions at the 190 m blind turn, Rung 3 succumbs to the standstill local optimum (95% stall rate). Multi-camera forward pruning (Rung 2) and multimodal LiDAR fusion (Rung 4) sharply eliminate stall failures and suppress total infractions.

the numerical convergence of the BEV autoencoder to 0.0046 MSE loss.

C. Ablation: VAE Architectural Capacity and Latent Bottleneck Dimensionality

To isolate whether the sensory scaling bottlenecks observed in Rung 2 and Rung 3 originate from insufficient autoencoder capacity or suboptimal bottleneck width, we conducted a controlled empirical ablation across network depth, convolutional filter width, and latent space dimensionality




Fig. 6. Route completion percentage across training timesteps in Town01. Rung 1 converges to 39.0% (195 m bottleneck); Rung 2 achieves rapid forward progress (25.0% in 12.2k steps); Rung 3 completely flatlines at 3.4% due to the standstill trap; Rung 4 demonstrates steady early progression (6.0% at 13.2k steps).





Fig. 7. Mean cumulative episode reward progression across training episodes. Rung 1 converges to 313.21 ± 38.45 over 200 episodes; Rung 2 reaches 158.31 ± 26.80 within 120 episodes (12.2k steps); Rung 3 collapses to 17.35 ± 3.12 in the standstill trap; Rung 4 improves from -8.81 to -5.02 ± 4.18 over 230 episodes (13.2k steps), though with negative mean reward throughout.

$d_z \in \{95, 190\}$ (Table III). All models were trained on 12,000 Cityscapes semantic frames and evaluated via linear ridge affordance decodability (R^2), participation-ratio effective rank, and forward-pass inference throughput.

1) *The Information Bottleneck Sweet Spot* ($d_z = 95$): As shown in Table III, the baseline 4-layer autoencoder achieves the highest overall affordance decodability (mean $R^2 = 0.8170$, lane offset $R^2 = 0.9713$, heading $R^2 = 0.8842$) with compact parameterization (13.7M weights). This strikes the optimal Information Bottleneck frontier: compressing 38,400 input pixels ($160 \times 80 \times 3$) into a 95-dimensional latent vector ($404 \times$ spatial compression) while preserving actionable geometric telemetry.

2) *Filter Width Scaling: Capacity Inflation vs. Affordance Dilution*: Doubling the convolutional channel widths ($64 \rightarrow 128 \rightarrow 256 \rightarrow 512$) quadruples the parameter footprint to



54.60M weights. While the effective rank increases from 4.77 to 8.29 (+73.8%), linear decodability of lane offset degrades significantly ($R^2 = 0.9713 \rightarrow 0.8821$). Because unregularized wide filters possess excessive capacity relative to the semantic input manifold, the network allocates latent variance toward high-frequency peripheral distractors (building textures, distant trees, sky gradients) rather than core driving affordances.

3) *Latent Space Scaling* ($d_z = 95 \rightarrow 190$): *The Invariance of Effective Rank*: Comparing **Deep-95** ($d_z = 95$) against **Deep-190** ($d_z = 190$) provides empirical evidence that the latent bottleneck is over-provisioned. When doubling the latent dimension from 95 to 190, the empirical effective rank remains virtually identical:

$$\text{Rank}_{\text{eff}}(d_z = 95) = 7.06 \quad \longleftrightarrow \quad \text{Rank}_{\text{eff}}(d_z = 190) = 7.14 \quad (24)$$

This suggests that the encoder concentrates meaningful variance into approximately 7–8 principal directions, consistent with the relatively low-dimensional driving affordances in CARLA’s urban scenes. Expanding d_z to 190 introduces ≈ 183 low-variance dimensions (reducing the effective utilization ratio from 7.4% down to 3.8%). We note that this observation is based on only two bottleneck widths within a single architecture and dataset; broader sweeps across d_z and additional environments would be needed to establish a more precise characterization of the latent manifold.

4) *Zero-Shot Policy Misalignment in Closed-Loop RL*: In closed-loop CARLA rollouts, hot-swapping newly trained VAE encoders into the pre-trained PPO policy induced an immediate 100% standstill stall rate (0.0 m route completion). This failure stems directly from latent coordinate basis rotation and activation under-excitation: while the baseline policy was trained on high-variance latents ($\sigma_z = 11.13$, max 42.78), the regularized Wide and Deep encoders maintain tightly bounded unit variance ($\sigma_z = 0.76$ and 0.68). Feeding $15\times$ smaller inputs into the first dense layer of the frozen actor suppressed longitudinal throttle activations below the braking threshold, driving the vehicle into an immediate stall trap. This confirms that representation capacity cannot be inferred from naive zero-shot policy reuse, but requires dedicated representation probing.

Detailed examination of simulation telemetry identified four primary failure mechanisms:

TABLE III
VAE ARCHITECTURAL CAPACITY, INFORMATION DECODABILITY, AND LATENT SPACE DIMENSIONALITY ABLATION

Configuration	Conv Depth	Filter Channel Hierarchy	Latent Dim (d_z)	Encoder Params	Eff. Rank $\uparrow$	Eff. Ratio (%)	Mean $R^2 \uparrow$	Lane Offset $R^2 \uparrow$	Heading $R^2 \uparrow$	Throughput (FPS) $\uparrow$
Baseline (Sweet Spot)	4	32 $\rightarrow$ 64 $\rightarrow$ 128 $\rightarrow$ 256	95	13,749,534	4.77	5.0%	0.8170	0.9713	0.8842	26.8 FPS (37.3 ms)
Wide (Over-Param.)	4	64 $\rightarrow$ 128 $\rightarrow$ 256 $\rightarrow$ 512	95	54,602,110	8.29	8.7%	0.7115	0.8821	0.7991	43.3 FPS (23.1 ms)
Deep (Hierarchical)	6	32 $\rightarrow$ 64 $\rightarrow$ 64 $\rightarrow$ 128 $\rightarrow$ 256 $\rightarrow$ 512	95	29,198,046	7.06	7.4%	0.6913	0.8965	0.7285	248.1 FPS (4.0 ms)
Deep-190 (Large Latent)	6	32 $\rightarrow$ 64 $\rightarrow$ 64 $\rightarrow$ 128 $\rightarrow$ 256 $\rightarrow$ 512	190	29,528,096	7.14	3.8%	0.7089	0.8847	0.7800	258.0 FPS (3.9 ms)

5) *Monocular Meter 190 Turn Blindness (Rung 1)*: As illustrated in Fig. 8, the Town01 benchmark transitions from a high-speed 190 m straightaway into an unbanked 90° right turn. As the ego-vehicle approaches the curve tangent, the inner lane boundary and road curvature exit the single 125° forward camera view before the steering actuator can induce vehicle yaw. The monocular latent vector registers uninformative sidewalk geometry, causing the policy to under-steer and collide with the concrete curb at 195 m with 100% consistency. In unseen Town02, zero-shot generalization degraded to $14.2\% \pm 3.1\%$ completion (Table II).

6) *Forward 3-Camera Rig Performance (Rung 2)*: By introducing angled cameras at $\pm 60^\circ$, Rung 2 provides lateral sightlines up to 210° . Despite the wider field of view, route completion degraded from the monocular baseline (39.0% to $25.0\% \pm 4.1\%$), with a maximum trajectory distance of 125 m—well short of the 190 m turn. The agent terminated via curb collision on the straight section (lane deviation 0.28 ± 0.08 m). The degradation may be attributable to the $2.9\times$ increase in state dimensionality (100-dim $\rightarrow$ 290-dim) combined with the substantially shorter training budget (12.2k steps vs. Rung 1’s ~ 105 k steps), making it difficult to disentangle the effect of camera count from training duration.

7) *Root-Cause Failure Mechanisms of 4-Camera Surround (Rung 3)*: In Rung 3, scaling from 3 to 4 cameras by adding a 180° rear camera caused catastrophic performance collapse ($3.4\% \pm 0.8\%$ completion, 17 m). This failure was driven by four interconnected mechanisms:

- 1) **High-Dimensional State Space Scaling**: Concatenating four 95-dim visual latents creates a 380-dimensional perceptual state space ($\mathbf{s}_t \in \mathbb{R}^{385}$). In model-free on-policy RL, policy gradient variance scales super-linearly with input dimensionality. Advantage gradients $\nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) \hat{A}_t$ become diluted across ungrounded peripheral pixels (e.g., empty sky, distant building tops), destabilizing policy updates.
- 2) **Observation Scale Mismatch**: The 380 latent dimensions have standard deviations $\sigma \approx 11$ –87 while the 5 navigation features are $O(1$ –20). Without observation normalization, the latent inputs dominate the first tanh layer and effectively drown the navigation telemetry. This scale mismatch grows with the number of concatenated latent vectors.
- 3) **The Standstill Local Optimum (Braking Trap)**: Because exploratory steering and throttle actions in high-dimensional state spaces frequently triggered immediate collisions or lane departures (penalized with $R_{\text{penalty}} = -10.0$), the policy gradient converged upon the trivial

local optimum: setting longitudinal throttle to 0 and applying continuous brake ($a_{\text{long}} \rightarrow -1.0$). By standing perfectly still at 17 m (mean lane deviation 0.05 ± 0.02 m), the agent avoids terminal crash penalties, accumulating a flat, non-negative reward ($R = 17.35 \pm 3.12$) until the simulation timeout. The absence of an idle penalty or stall termination condition in the reward design enables this exploit.

- 4) **Simulation Infrastructure Crash**: In Town02 transfer evaluations, maintaining four concurrent semantic segmentation camera viewports caused a Windows STATUS_STACK_BUFFER_OVERRUN (0xC0000409) crash in dense downtown intersections. Despite the reported VRAM usage of only 3150 MB on a 24 GB GPU, the crash may stem from Unreal Engine 4’s internal buffer management rather than raw GPU memory exhaustion.

8) *Closed-Loop Lane Departure in Multimodal VAEs (Rung 4)*: Across episodes 180–230 of Rung 4, pairing front vision with 2D BEV LiDAR completely eliminated the standstill trap, enabling smooth forward driving for 60–76 simulation steps (25–30 m) with positive rewards up to $+19.84$ (mean $R = -5.02 \pm 4.18$, lane deviation 0.35 ± 0.09 m). However, all trajectories terminated at the 3.0 m lane boundary after only ~ 30 m (6.0% route completion). Several factors likely contribute to this early termination: (i) the training budget of only 13.2k steps is substantially less than Rung 1’s ~ 105 k steps; (ii) the LiDAR BEV height clip of $[-2.0, +2.0]$ m relative to the sensor at $z = 2.4$ m may erase curb geometry (~ 0.15 m above ground), limiting the BEV channel’s ability to perceive road boundaries; and (iii) the latent variance mismatch between visual ($\sigma \approx 11$) and BEV ($\sigma \approx 87$) encoders, combined with the absence of observation normalization, may saturate the policy’s tanh activations. We note that the policy does receive ground-truth lane deviation d_{center} and heading error θ_{heading} in $\mathbf{x}_{\text{nav}}$, so the steering challenge is not purely perceptual but likely reflects insufficient training and the optimization difficulties of the high-variance multimodal input space.

VI. ARCHITECTURAL DISCUSSION: THE VAE SCALING CEILING AND THE MOTIVATION FOR ATTENTION

The empirical results of this ladder illustrate limitations encountered when scaling multi-sensor systems via monolithic autoencoders and flat latent concatenation.

A. The Structural Limitations of VAE Latent Concatenation

- 1) **Destruction of Spatial Inductive Bias**: Both perspective RGB and BEV LiDAR autoencoders map 2D spatial feature grids into 1D continuous latent vectors ($\mathbf{z} \in \mathbb{R}^{95}$)

via flattening and dense linear projection. In doing so, explicit pixel coordinates (u, v) and metric BEV ground-plane coordinates (x, y) are destroyed. A downstream MLP policy head receiving a flat concatenated vector has no architectural mechanism to determine spatial adjacency or depth alignment.

- 2) **Isolated Latent Compression (The Silo Effect):** In Rung 4, the visual VAE and LiDAR VAE compress inputs in complete isolation. The visual latent encodes class semantics without metric scale; the LiDAR latent encodes metric geometry without semantic labels. Simply concatenating them into $[\mathbf{z}_{\text{cam}} \parallel \mathbf{z}_{\text{lidar}}] \in \mathbb{R}^{190}$ forces a shallow MLP to discover complex cross-modal geometric correspondences (e.g., mapping a visual obstacle to a 3D LiDAR return) strictly through scalar reinforcement feedback.
- 3) **Gradient Dissipation Across Ungrounded Latents:** In high-dimensional flat latent spaces (290-dim in Rung 2, 385-dim in Rung 3), PPO advantage gradients are uniformly backpropagated across hundreds of continuous latent dimensions. Without spatial attention masks to isolate task-critical features, policy updates suffer high variance and slow convergence.

B. The Transition to Patch-Based Cross-Attention and VLA Transformers

These empirical observations motivate the investigation of patch-based tokenization and cross-attention transformer architectures [5], [7], [19] as a potential solution to these limitations:

- **Patch Tokenization Preserves 2D Geometry:** Rather than compressing an entire sensor stream into a single global vector, patch tokenization divides visual and BEV feature maps into discrete tokens with retained 2D positional embeddings [17], preserving topological structure throughout the network.
- **Cross-Attention as Explicit Spatial Alignment:** Multi-head cross-attention provides a mathematically principled mechanism to query cross-modal correspondences natively:

$$\text{Attn}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}_b^T}{\sqrt{d_k}}\right)\mathbf{V}_b \quad (25)$$

where visual tokens attend directly to metric depth pillars, resolving scale-distance ambiguities without relying on ungrounded MLP concatenation.

- **Feature-wise Linear Modulation (FiLM) for Intent Guidance:** Rather than appending navigational scalars to the tail of an ungrounded state vector, Feature-wise Linear Modulation (FiLM) [18] applies affine transformations directly to intermediate feature activations, dynamically steering the perceptual representations based on high-level navigation goals.

VII. CONCLUSION

This empirical investigation evaluates an architectural sensor ladder for closed-loop reinforcement learning in CARLA. The

monocular baseline achieves the highest route completion (39.0%) after $\sim 105\text{k}$ training steps but fails at the 90° turn at Meter 190. Adding cameras (Rungs 2–3) degrades performance, with the 4-camera configuration collapsing into a standstill trap exploiting the absence of an idle penalty. Camera-LiDAR fusion (Rung 4) eliminates stalling but achieves only 6.0% completion at 13.2k steps. We observe that flat VAE latent concatenation encounters representational and optimization difficulties as sensor count increases, compounded by observation scale mismatches and substantially shorter training budgets for the multi-sensor configurations. These observations motivate future work on patch-based cross-attention architectures, equal-budget comparisons with multiple seeds, and observation normalization to better isolate the effect of sensor scaling from optimization confounds.

REFERENCES

- [1] D. A. Pomerleau, “ALVINN: An autonomous land vehicle in a neural network,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, 1988, pp. 305–313.
- [2] M. Bojarski *et al.*, “End to end learning for self-driving cars,” *arXiv preprint arXiv:1604.07316*, 2016.
- [3] F. Codevilla, M. Müller, A. López, V. Koltun, and A. Dosovitskiy, “End-to-end driving via conditional imitation learning,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2018, pp. 4693–4700.
- [4] A. Dosovitskiy, G. Ros, F. Codevilla, A. Lopez, and V. Koltun, “CARLA: An open urban driving simulator,” in *Proc. Conf. Robot Learn. (CoRL)*, 2017, pp. 1–16.
- [5] A. Prakash, K. Chitta, and A. Geiger, “Multi-modal fusion transformer for end-to-end autonomous driving,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2021, pp. 7077–7087.
- [6] D. Chen, B. Zhou, V. Koltun, and P. Krähenbühl, “Learning by cheating,” in *Proc. Conf. Robot Learn. (CoRL)*, 2020, pp. 66–75.
- [7] H. Shao, L. Wang, R. Chen, H. Li, and Y. Liu, “Safety-enhanced autonomous driving using interpretable sensor fusion transformer,” in *Proc. Conf. Robot Learn. (CoRL)*, 2022, pp. 726–737.
- [8] A. H. Lang, S. Vora, H. Caesar, L. Zhou, J. Yang, and O. Beijbom, “PointPillars: Fast encoders for object detection from point clouds,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2019, pp. 12697–12705.
- [9] J. Huang, G. Huang, Z. Zhu, Y. Ye, and D. Du, “BEVDet: High-performance multi-camera 3D object detection in bird-eye-view,” *arXiv preprint arXiv:2112.11790*, 2022.
- [10] M. Toromanoff, E. Wirbel, and F. Moutarde, “End-to-end model-free reinforcement learning for urban driving using implicit affordances,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2020, pp. 7153–7159.
- [11] Z. Zhang, A. Liniger, R. Sakai, and F. Yu, “End-to-end autonomous driving with Roach,” in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2021, pp. 2875–2885.
- [12] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” *arXiv preprint arXiv:1707.06347*, 2017.
- [13] J. Schulman, P. Moritz, S. Levine, M. I. Jordan, and P. Abbeel, “High-dimensional continuous control using generalized advantage estimation,” in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2016.
- [14] S. Mysore, B. Mabsout, R. Mancuso, and K. Saenko, “Regularizing action policies for smooth control with reinforcement learning,” in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2021, pp. 1810–1816.
- [15] D. P. Kingma and M. Welling, “Auto-encoding variational Bayes,” in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2014.
- [16] M. Cordts *et al.*, “The Cityscapes dataset for semantic urban scene understanding,” in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 3213–3223.
- [17] A. Dosovitskiy *et al.*, “An image is worth 16x16 words: Transformers for image recognition at scale,” in *Proc. Int. Conf. Learn. Represent. (ICLR)*, 2021.
- [18] E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville, “FiLM: Visual reasoning with a general conditioning layer,” in *Proc. AAAI Conf. Artif. Intell.*, 2018, pp. 3942–3950.

- [19] A. Vaswani *et al.*, “Attention is all you need,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2017, pp. 5998–6008.